Two camps break the consistent core: set-theoretic estimation over 159 ratification documents

Federalist + Anti-Federalist, scored by the proved STE algorithm

thejeb project note

2026-07-28

Abstract

The single-camp Federalist run (86 essays) found a large *consistent core*: 7 of 10 variables unanimous where spoken, only 3 conflict sites. We test whether the STE disagreement structure tracks corpus composition by adding the opposing camp — 73 **Anti-Federalist** essays (Brutus, Cato, Centinel, Federal Farmer, An Old Whig, Philadelphiensis; public domain) — for a **159-document two-camp corpus**. The same chunking pipeline extracts a canonical 10-variable schema (now carrying both poles of each question) and one frame per document; the verified `tides` executable scores all 159 voices with `STE.Consistent` and `STE.disagreementDegree` (`Ste.FiniteInstance`). The verdict: the feasibility set is again empty, but where the Federalist corpus had **3** conflict sites the two-camp corpus has **all 10** — every variable is contested, and each split runs cleanly along the Federalist/Anti-Federalist line. The disagreement profile is a function of the corpus, computed by the proved algorithm, not a fixed property of the topic.

1 The experiment

The Federalist companion note found that 86 essays by Publius share a large consistent core. The sharp test of whether that structure is intrinsic or compositional is to add the essays arguing the other side. We fetch 73 Anti-Federalist essays from public-domain sources (`scripts/antifederalist/`) and pool them with the 86 Federalist essays into one corpus of **159 documents**, each a single STE voice. The pipeline (`scripts/corpus/`, now accepting multiple corpus directories) runs unchanged: pass 1 (Haiku, per chunk) over 262 chunks proposes 3,253 open (variable, value) pairs; pass 2 (Sonnet) unifies them into a canonical **10-variable** schema whose domains now include *both* positions on each contested question; pass 3 (Haiku, per document) emits one frame per essay. The `frames.json` is scored by the verified executable, which evaluates the definitions proved in `Ste.FiniteInstance` on the loaded data (no `native_decide`).

2 The verdict: the core shatters

`STE.Consistent` = false: no single reading satisfies all 159 voices. But the disagreement is now *total*. Every one of the 10 variables is a conflict site (disagreement degree ≥ 2), against 3 in the Federalist-only run:

variable	Federalist-only	two-camp
ratification_stance	1	**3**
governance_structure	3	3
standing_army_peacetime	1	**2**
bill_of_rights_necessity	1	**2**
federal_taxation_scope	—	**2**
executive_power_character	1	**2**
judiciary_power_scope	1	**2**
house_representation_adequacy	—	**2**
confederation_adequacy	1	**2**
senate_character	2	**2**
conflict sites	**3**	**10**

(The two schema runs name variables slightly differently; the invariant claim is the count of conflict sites, $3 \to 10$.)

3 The split is clean and camp-driven

The degrees are not extraction noise: spot-checks against the frames show each conflict is the Federalist/Anti-Federalist divide.

variable	Federalist essays	Anti-Federalist essays
ratification_stance	favor (78)	oppose (55), amend (12)
bill_of_rights_necessity	unnecessary (2)	necessary (25)
standing_army_peacetime	safe (10), dangerous (2)	dangerous (27), safe (1)

The degree-3 ratification_stance split is exactly the two camps plus the conditional-ratification wing (favor_with_amendments). The bill_of_rights axis is cleanly opposite (Federalist 84 vs. the Anti-Federalist demand). The standing_army split is near-perfect, the few crossovers reflecting genuine nuance in the sources.

4 Three shapes, one algorithm

The same verified `Consistent` / `disagreementDegree` now records three qualitatively different disagreement shapes:

corpus	consistent?	structure
Tides (7)	no	pervasive, multi-camp; worst degree 5
Federalist (86)	no	large consistent core; 3 conflict sites
Ratification (159)	no	total shatter; all 10 variables contested

Adding the opposing 73 essays flips seven unanimous variables into conflicts. This is precisely the behavior of the modal STE theory: acquiring information shrinks the feasible family; here it empties the agreement on *every* axis at once. The disagreement profile is compositional and computed, not a story told about the topic.

5 The boundary of trust

The verdict is the verified STE definitions evaluated on runtime-parsed frames; each asserted value in `frames.json` carries a licensing quote, so every cell is auditable against the essay, and the three headline splits were checked against those quotes. A mis-extraction would move a count, but the shatter — ten contested variables where the single camp had three — is far larger than any plausible extraction slack.

References